

# Project Deliverable 4

## Group 6:

Rimjhim Jain  
Sudhanshu Dalvi  
Saumit Chinchkhadi  
Naval Gaikwad  
Aniket Ghumed

## AWS Cost Analysis for Amazon Best Selling Prediction Project

### AWS Services Utilized are:

- **AWS S3 (Simple Storage Service):**

Purpose: Used for storing datasets.

Cost Factors: Storage capacity, requests, and data transfer.

- **Amazon Athena:**

Purpose: Querying transformed data.

Cost Factors: Based on the amount of data scanned by the queries.

- **AWS Glue:**

Purpose: ETL (Extract, Transform, Load) jobs for data transformation.

Cost Factors: Number of Data Processing Units (DPUs) used and job run time.

- **Amazon QuickSight:**

Purpose: Creating visualizations for data insights.

Cost Factors: Monthly subscription and additional charges based on user count and SPICE (Super-fast, Parallel, In-memory Calculation Engine) capacity.

- **Amazon SageMaker:**

Purpose: Model training, evaluation, and hyperparameter tuning.

Cost Factors: Instance hours, data processing, and storage.

- **AWS Lambda (optional for scaling):**

Purpose: Automating and scaling ETL jobs.

Cost Factors: Number of requests and duration of code execution.

## Cost and usage [Info](#)

Current month costs

**\$1.92**

Forecasted month end costs

There isn't enough historical data to forecast your spend

Last month costs


**\$1.16**

Average month costs

**\$0.51**

Total costs per month

Cost (USD)



SageMaker S3 CloudWatch  
Glue Athena

[Go to Billing and Cost Management](#)



## Cost summary [Info](#)

Month-to-date cost

**\$1.92**

- compared to last month for same period

Last month's cost for same time period

**\$0.00**

Nov 1 – 10

Total forecasted cost for current month

Access denied

Last month's total cost

**\$1.16**

let's analyze the AWS costs for your Cloud-Based Machine Learning Project:

#### Month-to-Date Cost Analysis

Month-to-Date Cost: \$1.92

This is the total cost incurred from the start of the current month up to the present day.

It shows an increase compared to the last month for the same period, where the cost was \$0.00. This indicates new usage or a scaling up of services.


Last Month's Total Cost: \$1.16

The total cost for the previous month was \$1.16, which is less than the current month-to-date cost. This suggests either an expansion in the project's scope or increased usage of AWS services.

#### Cost Breakdown by Service

Amazon SageMaker, Amazon Simple Storage Service (S3), AWS Glue, etc.

The cost breakdown by service is useful for understanding which AWS services are contributing most to the overall cost.

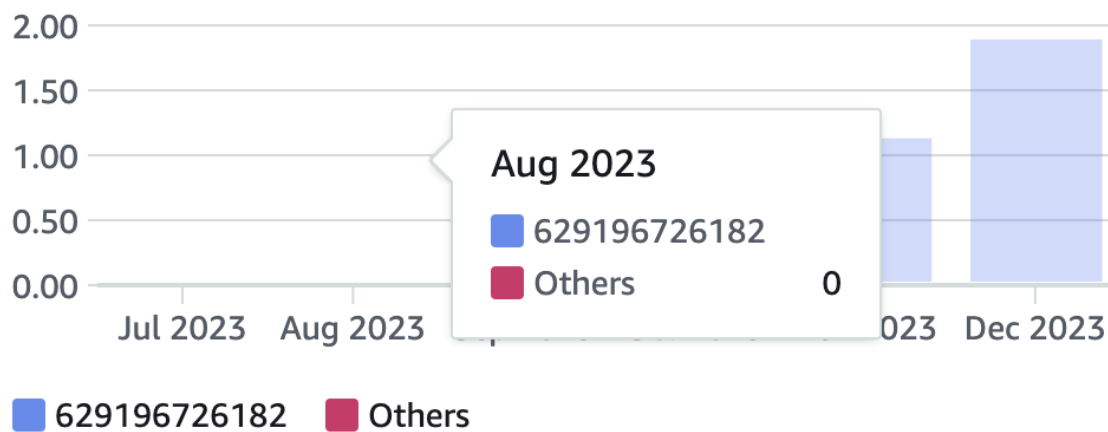


## ⚙️ Cost breakdown [Info](#)

Group costs by

Account ▼

Costs (\$)



[Analyze your costs in Cost Explorer](#)



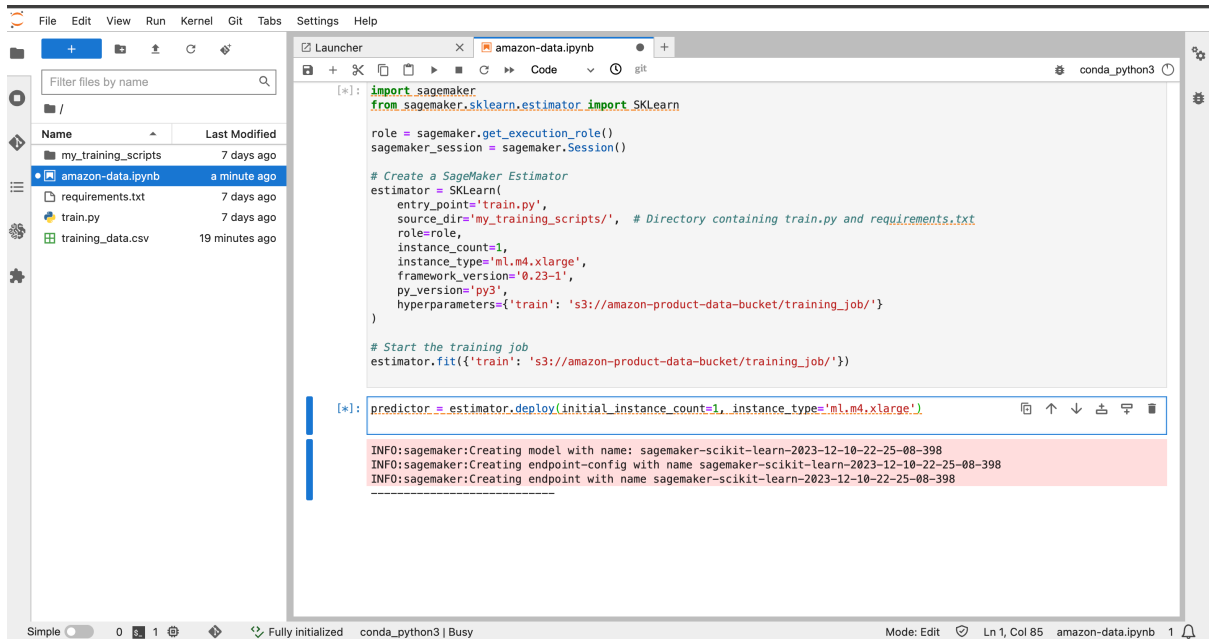
### Cost Optimization Strategies:

- Regularly monitor and review AWS usage and costs using AWS Cost Explorer.
- Optimize data storage by cleaning up unused S3 objects.
- Manage Athena query costs by optimizing query efficiency.
- Utilize SageMaker spot instances for cost-effective model training.
- Consider AWS reserved instances for predictable workloads.

### Conclusion after cost analysis:

Effective cost management is crucial for the sustainability of the project. Regular monitoring and optimization strategies should be implemented to ensure that the project remains within budget while achieving its objectives.

# Deployment:



The screenshot shows a Jupyter Notebook titled 'amazon-data.ipynb' in a web-based IDE. The left sidebar displays a file explorer with a search bar and a list of files: 'my\_training\_scripts' (7 days ago), 'amazon-data.ipynb' (a minute ago), 'requirements.txt' (7 days ago), 'train.py' (7 days ago), and 'training\_data.csv' (19 minutes ago). The main area contains Python code for creating a SageMaker estimator and deploying it. The code includes imports for 'sagemaker', 'sagemaker.sklearn.estimator', and 'SKLearn'. It defines a role, creates a SageMaker estimator with specific instance types and hyperparameters, and then deploys it. The output shows the deployment process, including the creation of a model and an endpoint configuration.

```
[*]: import sagemaker
from sagemaker.sklearn.estimator import SKLearn

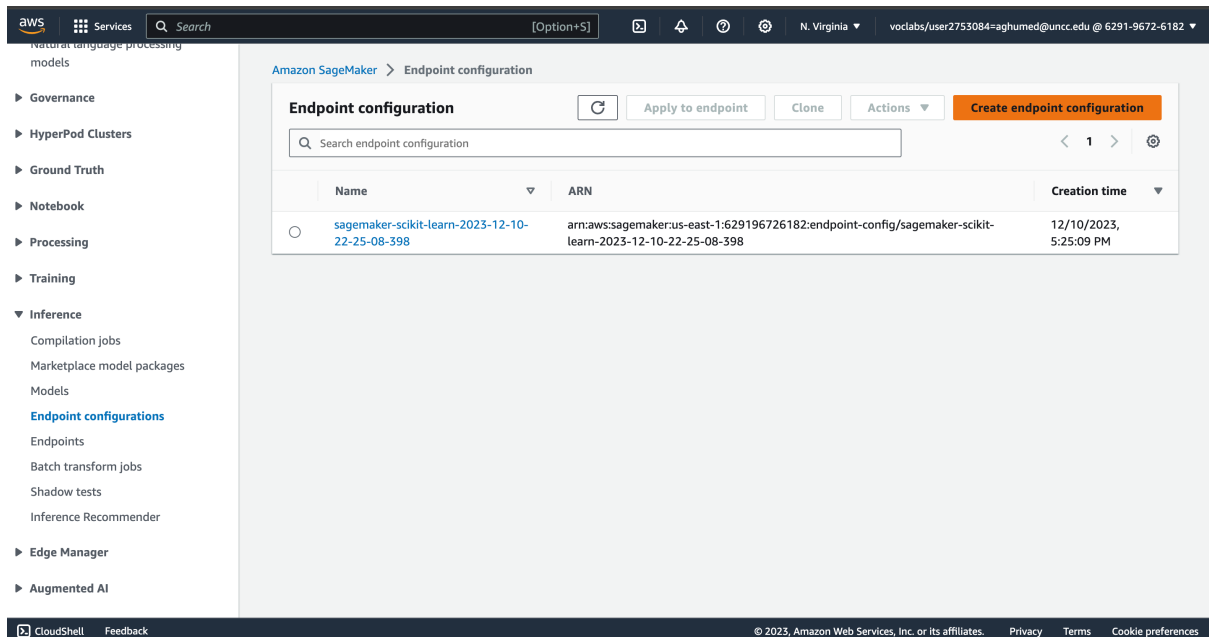
role = sagemaker.get_execution_role()
sagemaker_session = sagemaker.Session()

# Create a SageMaker Estimator
estimator = SKLearn(
    entry_point='train.py',
    source_dir='my_training_scripts/', # Directory containing train.py and requirements.txt
    role=role,
    instance_count=1,
    instance_type='ml.m4.xlarge',
    framework_version='0.23-1',
    py_version='py3',
    hyperparameters={'train': 's3://amazon-product-data-bucket/training_job/'})

# Start the training job
estimator.fit({'train': 's3://amazon-product-data-bucket/training_job/'})

[*]: predictor = estimator.deploy(initial_instance_count=1, instance_type='ml.m4.xlarge')
```

INFO:sagemaker:Creating model with name: sagemaker-scikit-learn-2023-12-10-22-25-08-398  
INFO:sagemaker:Creating endpoint-config with name sagemaker-scikit-learn-2023-12-10-22-25-08-398  
INFO:sagemaker:Creating endpoint with name sagemaker-scikit-learn-2023-12-10-22-25-08-398



The screenshot shows the Amazon SageMaker console interface. The left sidebar contains a navigation menu with categories like 'Governance', 'HyperPod Clusters', 'Ground Truth', 'Notebook', 'Processing', 'Training', 'Inference', 'Edge Manager', and 'Augmented AI'. The 'Inference' section is expanded, showing options like 'Compilation jobs', 'Marketplace model packages', 'Models', 'Endpoint configurations', 'Endpoints', 'Batch transform jobs', 'Shadow tests', and 'Inference Recommender'. The main area displays the 'Endpoint configuration' page. It includes a search bar, a table of configurations, and a 'Create endpoint configuration' button. The table shows a single configuration with the name 'sagemaker-scikit-learn-2023-12-10-22-25-08-398' and an ARN 'arn:aws:sagemaker:us-east-1:629196726182:endpoint-config/sagemaker-scikit-learn-2023-12-10-22-25-08-398'. The creation time is '12/10/2023, 5:25:09 PM'.

Amazon SageMaker > Endpoint configuration

Endpoint configuration

Search endpoint configuration

| Name   | ARN   | Creation time          |
|--|---|------------------------|
| sagemaker-scikit-learn-2023-12-10-22-25-08-398 | arn:aws:sagemaker:us-east-1:629196726182:endpoint-config/sagemaker-scikit-learn-2023-12-10-22-25-08-398 | 12/10/2023, 5:25:09 PM |

Natural language processing models

► Governance

► HyperPod Clusters

► Ground Truth

► Notebook

► Processing

► Training

▼ Inference

Compilation jobs

Marketplace model packages

Models

Endpoint configurations

Endpoints

Batch transform jobs

Shadow tests

Inference Recommender

► Edge Manager

► Augmented AI

Amazon SageMaker > Endpoints

Endpoints

↻

Update endpoint

Actions ▼

Create endpoint

🔍 Search endpoints

< 1 >

|   | Name ▼   | ARN ▼  | Creation time ▼        | Status ▼   | Last updated ▼         |
|---|--|--|------------------------|------------|------------------------|
| ○ | <a href="#">sagemaker-scikit-learn-2023-12-10-22-25-08-398</a> | arn:aws:sagemaker:us-east-1:629196726182:endpoint/sagemaker-scikit-learn-2023-12-10-22-25-08-398 | 12/10/2023, 5:25:09 PM | 🕒 Creating | 12/10/2023, 5:25:09 PM |

CloudShell Feedback

© 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Getting started

Studio

Studio Lab

Canvas

RStudio

TensorBoard

Profiler

▼ Admin configurations

Domains

Role manager

Images

Lifecycle configurations

SageMaker dashboard

Search

▼ JumpStart

Foundation models

Computer vision models

Amazon SageMaker > Models

Models

↻

Create endpoint

Create endpoint configuration

Actions ▼

Create model

🔍 Search models

< 1 > ⚙️

|   | Name ▼   | ARN   | Creation time ▼        |
|---|--|---|------------------------|
| ○ | <a href="#">sagemaker-scikit-learn-2023-12-10-22-25-08-398</a> | arn:aws:sagemaker:us-east-1:629196726182:model/sagemaker-scikit-learn-2023-12-10-22-25-08-398 | 12/10/2023, 5:25:08 PM |

CloudShell Feedback

© 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Services

Search

[Option+S]

N. Virginia

voclabs/user2753084=aghumed@uncc.edu @ 6291-9672-6182

Notebook

Processing

Training

Inference

Compilation jobs

Marketplace model packages

Models

Endpoint configurations

Endpoints

Batch transform jobs

Shadow tests

Inference Recommender

Edge Manager

Augmented AI

AWS Marketplace

Tutorials

Documentation

Amazon SageMaker > Endpoints > sagemaker-scikit-learn-2023-12-11-00-03-30-518

sagemaker-scikit-learn-2023-12-11-00-03-30-518

Delete

Endpoint summary

|      |  |                      |   |              |   |
|------|--|----------------------|---|--------------|---|
| Name | sagemaker-scikit-learn-2023-12-11-00-03-30-518   | Status               | InService   | Type         | Real-time   |
| ARN  | arn:aws:sagemaker:us-east-1:629196726182:endpoint/sagemaker-scikit-learn-2023-12-11-00-03-30-518                       | Creation time        | Sun Dec 10 2023 19:03:31 GMT-0500 (Eastern Standard Time)               | Last updated | Sun Dec 10 2023 19:06:54 GMT-0500 (Eastern Standard Time) |
| URL  | https://runtime.sagemaker.us-east-1.amazonaws.com/endpoints/sagemaker-scikit-learn-2023-12-11-00-03-30-518/invocations | Model container logs | /aws/sagemaker/endpoints/sagemaker-scikit-learn-2023-12-11-00-03-30-518 | Alarms       | 0 alarms  |

Monitor

Settings

Alarms

CloudShell

Feedback

© 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

File Edit View Run Kernel Git Tabs Settings Help

Filter files by name

/ my\_training\_scripts /

| Name             | Last Modified |
|------------------|---------------|
| requirements.txt | 7 days ago    |
| train.py         | an hour ago   |

amazon-data.ipynb

train.py

train.py

Code

conda\_python3

```
SM_CHANNEL_TRAIN=/opt/ml/input/data/train
SM_HP_TRAIN=s3://amazon-product-data-bucket/training_job/
PYTHONPATH=/opt/ml/code:/miniconda3/bin:/miniconda3/lib/python3.7:/miniconda3/lib/python3.7/lib-dynload:/miniconda3/lib/python3.7/site-packages
Invoking script with the following command:
/miniconda3/bin/python train.py --train s3://amazon-product-data-bucket/training_job/
Requirement already satisfied: fsspec in /miniconda3/lib/python3.7/site-packages (2023.1.0)
WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is recommended to use a virtual environment instead: https://pip.pypa.io/warnings/venv
[notice] A new release of pip is available: 23.0 -> 23.3.1
[notice] To update, run: pip install --upgrade pip
sys:1: DtypeWarning: Columns (13) have mixed types.Specify dtype option on import or set low_memory=False.

2023-12-10 23:56:03 Uploading - Uploading generated training model
precision recall f1-score support
False 0.99 1.00 1.00 283544
True 0.33 0.01 0.01 1724
accuracy 0.99 283268
macro avg 0.66 0.50 0.51 285268
weighted avg 0.99 0.99 0.99 285268
2023-12-10 23:55:58,442 sagemaker-containers INFO Reporting training SUCCESS

2023-12-10 23:56:19 Completed - Training job completed
Training seconds: 177
Billable seconds: 177

[38]: predictor = estimator.deploy(initial_instance_count=1, instance_type='ml.m4.xlarge')

INFO:sagemaker:Creating model with name: sagemaker-scikit-learn-2023-12-11-00-03-30-518
INFO:sagemaker:Creating endpoint-config with name sagemaker-scikit-learn-2023-12-11-00-03-30-518
INFO:sagemaker:Creating endpoint with name sagemaker-scikit-learn-2023-12-11-00-03-30-518
-----!
```

Simple

0 1

Fully initialized conda\_python3 | Idle

Mode: Edit Ln 2, Col 1 amazon-data.ipynb 1